

“But”: Pseudo-Interactive AIGC Film

Exploring the Illusion of Agency through Simulated Interaction and Game-Style UI

Shiqin Hou

HKUST

shouaf@connect.ust.hk

Abstract

“But” is a 6-minute AI-generated narrative short that treats the illusion of choice as its central subject. Within an integrated AI platform (Tapnow), a persistent on-screen cursor and three game-style notification interfaces frame the film as a predetermined “system session”: the viewer is invited to act, yet every branch is authored in advance. This technical report documents the end-to-end pipeline realized through a pseudo-interactive AIGC filmmaking framework (Gemini for scripting and stepwise shot planning, Midjourney v7 for scene stylization, Nano Banana Pro for character and UI assets under a unified aesthetic, Seedance 2.0 for animation, CapCut for compositing), the methods used to hold a single protagonist coherent across six costume changes, and the design of the pseudo-interactive UI as a formal shell for the film’s argument about false agency. This report discusses the work as a case of human-led, tool-mediated authorship in AIGC filmmaking.

CCS Concepts

• **Computing methodologies** → **Computer graphics**; • **Applied computing** → *Media arts*.

Keywords

AI-generated film; pseudo-interactive; game UI integration; character consistency; AIGC workflow; human-AI co-creation

ACM Reference Format:

Shiqin Hou. 2026. “But”: Pseudo-Interactive AIGC Film: Exploring the Illusion of Agency through Simulated Interaction and Game-Style UI. In *SIGGRAPH Asia 2026 AI Film Frontiers Workshop (SA '26)*. ACM, New York, NY, USA, 4 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

But is a 6-minute AIGC pseudo-interactive narrative short. This report proposes a pseudo-interactive AIGC filmmaking framework – a single authoring method that orchestrates off-the-shelf generative tools into one pre-authored, world-model-style film space wrapped in a game-style interface. The framework and its argument are causally linked: by staging interaction as simulation, the work turns the viewer’s sense of control into a media-art claim about living inside systems that feel free yet hand us only the appearance

of choice. The work grew out of a lived contradiction of wage labor rather than a stylistic brief. Amid endless busy work, daily life was buried and turned rigid; people gradually became unthinking beings who could only mechanically perform the so-called tasks of living. Yet the moment I passed the pantry room and caught sight of that stretch of azure sky and white clouds, it all shattered in an instant. The gloom was gone, replaced by a clear blue sky. In that instant the inspiration for this short film took shape – a fresh awareness of observing the self rather than experiencing it firsthand. Surveying prior AI-film research shaped what this film sets out to do: Atzmon et al. [1] name the identity–motion trade-off that makes a consistent protagonist hard, and Zheng et al. [2] propose step-by-step, minimal-intervention generation – together they point to consistency and planning as the real engineering problems, not the models; Li et al. [5] and Shi et al. [6] orchestrate tools or agents for cross-scene consistency, while Huang et al. [3] show how human–LLM co-creation turns real footage into a fabricated narrative, a tension that AI film keeps returning to. Building on these, the film’s goal is to use AI-generated film as a legitimate artistic medium for a sober recognition – that accepting limitation may be the only real freedom – and to take the one move none of them make: where Chen et al. [4] drive a cinematic VR narrative through real-time, input-triggered generative-AI interaction, *But* inverts that premise and simulates interaction from a fixed script, turning the viewer’s sense of control into the work’s media-art claim. Carrying this out meant working entirely within commodity tools – Gemini for scripting, Nano Banana Pro for locked character assets, Midjourney v7 for real-garment fusion, Seedance 2.0 for animation, and CapCut for the interface shell (Table 1) – yet two difficulties arose: keeping one protagonist legible across six costumes and discontinuous time, and unfolding the pseudo-interaction so the viewer’s sense of control reads as real before it is withdrawn. The response was to build a small system rather than patch scenes: a reusable character-asset library, and a structure that withholds the film’s fantastical nature until the final moments, so the viewer first inhabits the daydream as real and is then retroactively shown it was simulation. The result is a 6-minute MP4 video (H.264, 2560×1440, 30 fps, Cantonese with English subtitles) whose recognition is not spoken aloud but delivered as a structure the viewer is made to feel.

2 The Pseudo-Interactive AIGC Framework

The generative stages of the pipeline run within **Tapnow**, an integrated AI creation platform that consolidates the model tools – Gemini for scripting, Midjourney v7 and Nano Banana Pro for image generation, and Seedance 2.0 for video clips – into one workspace; the generated clips and UI images are then assembled in CapCut, a separate desktop editor. The six design stages map onto these tools as Table 1 shows:

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SA '26, Kuala Lumpur, Malaysia

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-X-XXXXXX-XX-X/26/12

<https://doi.org/XXXXXXX.XXXXXXX>

Table 1: The framework instantiated across tools: each of the six design stages with its authorial method and realizing tool, plus the video engine (Seedance 2.0) used for scene motion.

Design stage	Authorial method	Tool
Scene-first	Script + shot planning	Gemini
Character assets	Locked face, 3-view sheet	Nano Banana Pro
Real-garment	Real outfits + style (real refs)	Midjourney v7 + real refs
Fragmentation	Intentional fractures	authorial
Multi-view	Base → multi-angle	Nano Banana Pro
UI shell	UI generated, then cut into the final edit	Nano Banana Pro + CapCut
Scene motion	Animates each generated scene	Seedance 2.0

The framework is a single production chain in which six design stages — (1) scene-first authoring, (2) a reusable character-asset library, (3) real-garment fusion, (4) deliberate cross-scene fragmentation, (5) base-shot multi-view decomposition, and (6) a game-style pseudo-interactive shell — are orchestrated across commodity tools (Table 1). Subsections 2.1–2.3 realize design stages (2)–(5); the shell (stage 6) is the subject of Subsection 3.1.

2.1 Multi-Appearance Character Consistency

Design stage (2): creating reusable assets and building a character-asset library.

The first major difficulty is keeping one recognizable protagonist stable across six costumes, six lighting setups, and six scenes. Text-to-video generators let a single prompt drift freely, so the same face arrives at different generations with different features, proportions, or styling — the identity–motion trade-off Atzmon et al. [1] describe, and the very reason cross-shot consistency is hard. Research systems meet it with orchestration: Anim-Director [5] sets a large multimodal model to work as an autonomous director-agent holding scenes coherent through visual-language prompting, while AniMaker [6] spreads the work across agents that keep a protagonist stable shot to shot. Both are bespoke, trained systems built for the research setting, not the off-the-shelf commercial tools used here; this film borrows only the paradigm — that consistency must be engineered rather than left to the generator — and removes the drift at the source instead of training a model to maintain consistency. The face was first produced by prompting yet was initially too neutral; working from that sample, the features were fine-tuned through further prompts and then locked as a three-view sheet with varied expressions that becomes a reusable asset fed into costume design, scene fusion, and video generation. Costumes, likewise, are not invented from text: real clothing references are fused with the film’s style keywords, replacing an earlier pure-prompt attempt whose texture and material quality fell short. The result is that audiences identify the same person through every costume change, which lets the final reveal — all outfits displayed at once — achieve its maximum impact.

The deliberate chromatic, costume, and temporal fractures between scenes are not defects of this pipeline but an authorial choice: each scene was generated independently, and the resulting discontinuity carries the film’s theme of a pluralistic life (Subsection 2.3).

The recurrence of one face is also what lets the film read as a single person’s routine rather than six unrelated figures. Because the same features survive every costume and lighting change, the viewer assembles one continuous subject out of clips that were never shot together, and that assembly is the work’s quiet claim about routine: the week repeats a protagonist, not a story. The sameness is authored, not observed, which is precisely the point — a stable, repeatable persona is delivered to the audience as given, and the recognition it triggers is recognition of a pattern the director fixed in advance.

2.2 Multi-Camera Spatial Continuity

Design stage (5): from a single plane to multiple three-dimensional views.

Creating the subway-station scene raised a spatial-continuity problem, and the added pseudo-interactive element made that problem more acute. The design required a first-person POV shot in which the camera appears to search for the protagonist; after she walks off-frame, two empty shots of the station play under electromagnetic-interference sound, a “FORCE EXIT” warning popup explains the interruption, and once the cursor is clicked the camera returns at higher speed, swings left–right, and finally finds her leaning against a wall, asking “Are you looking for me?”. Holding this POV insertion continuous with the surrounding special-angle shots — foreground overhead, side tracking, low-angle rotation — demands that every angle describe the *same* space. Rather than draw a 2D floor plan, the method anchors all angles to one generated base: a single-shot image of the intended location is produced first, then decomposed through prompting into multiple camera angles of that same spot, so geometry and object positions stay coherent across the cut. The mouse cursor, appearing inside the POV segment and persisting across the empty shots and the return, gives viewers a stable reference that also marks the screen-mediated nature of the experience. The interference audio and the “external interrupt detected” popup motivate the two empty shots, and a diegetic sound field — the rush of a passing train and the electric leakage of the glitch — implies off-screen movement, letting the audience reconstruct the station space even while the protagonist is absent.

Anchoring every angle to a single generated base keeps geometry coherent, but it also means the ‘multiple camera views’ are recombinations of one authored image rather than independent captures. That constraint is the point rather than a limitation: the sequence presents the illusion of a navigable space — a camera that searches, lingers, and returns — while every vantage is in fact derived from the same fixed frame. The coherence the viewer reads as spatial continuity is therefore pre-authored closure, the same closure that lets the empty POV shots read as one continuous station rather than a montage, and that the film’s world-model reading names as a closed, non-responsive system.

2.3 Fantasy and Reality Controllability

Design stage (3) and (4): fusing real garments and fragmenting scenes on purpose.

The contradictory narrative demand is that audiences should take the fantasy sequences as real until the final reveal, yet on

rewatching should catch the subtle clues that expose the artifice – a balance between immersion and foreshadowing that is fundamentally a problem of visual-style control. The film meets this demand through deliberate per-scene color design rather than a binary color code. Each scene carries a consciously chosen dominant tone – the convenience store’s black-and-white checkerboard floor at the entrance and pinkish tone inside, the bedroom’s cool-and-warm mix, the restaurant’s yellow, the stairwell’s cement gray, the subway station’s red (both in the malfunction and while waiting), the rooftop’s night-sky blue, and the office’s deep blue – yet none of these tones is assigned to mark a scene as fantasy or as real. The fantasy/reality boundary is instead carried by the protagonist’s seemingly arbitrary, narratively unmotivated costume changes, which blur imagined and real and suggest every scene belongs to one unstable mental space, and by scene transitions that jump illogically between day and night without causal motivation: a temporal disorientation, not the color scheme, that gives the experience its dreamlike, constructed feel and a deliberate sense of pseudo-chaos. Two “system loading” moments, one inside the convenience-store scene and one as the protagonist waits on the platform, use brief glitch effects and screen distortion that appear at first to be stylistic, then reveal themselves as diegetic hints about the medium. The three game interfaces (Subsection 3.1) add a further artificial-signal layer, reminding audiences they are watching a system or session and retrospectively reframing the whole fantasy portion as something occurring inside a digital frame.

2.4 Sound, Scene Order, and Passivity

Design stage (6): the pseudo-interactive UI shell and the post-production layer that binds the work.

Editing and assembly are the final creative act: they make the pseudo-interactive reading legible. The sequence of scenes, the timing of each glitch, the instant a pop-up appears on the screen – these determine how the work is read: the viewer is watching a session, not a story being told, thus the boundary between the fictional world and its medium remains clearly visible. This is why design stage (6) refers to the post-production layer: the argument doesn’t reside in any single shot, but in the blank spaces chosen during editing.

Two other production decisions, in terms of sound direction and narrative pacing, also reinforce the pseudo-interactive argument. The narrative soundscape, deliberately added in post-production, reflects a denial of the viewer’s autonomy, much like the visual style. Every scene in the film is laid over a bed of white noise and city-road traffic – cars passing in the distance – so the world always sounds urban and inhabited. This is further reinforced by the sound of a speeding train in the subway scene and the hum of leaking electricity in the glitch section; the final scene back in the office adds the atmosphere of a busy workplace and the ringing of a telephone, suggesting the protagonist is just one of many people repeating the same routine day after day. But none of these soundscapes respond to cursor clicks. These elements don’t become denser, sparser, or shifted by the triggering of “choices.”

The sequence of scenes – convenience store, bedroom, restaurant, stairwell, subway, rooftop birthday party, office – is cumulative, not progressive. The film doesn’t culminate in a climax; it

Table 2: The three pseudo-interactive interfaces.

#	Trigger	Accent	Function
1	Opening convenience-store scene	white/cyan	Session start
2	Subway scene (protagonist off-frame)	red	Warning message
3	Post-reveal hard cut	white/green	Statistics

loops through a week’s worth of work, the rooftop birthday party being just one segment in the loop, not a release. The form is pseudo-interactive, but the routine is real, and the sound and scene sequence subtly emphasize the gap between the two. The “FORCE EXIT” popups, system loading failures, and the ever-present mouse cursor throughout the editing appear as the final layer of this emphasis: the interface symbolizes autonomy, while the world rejects it, ultimately leaving the audience to confront both simultaneously. The final film, with all its materials – images, colors, sounds, scene sequence, interface – consistently expresses the same false agency claim, and it is this consistency that makes the pseudo-interactive form feel like a coherent stance, rather than a flaw in a workflow.

3 Pseudo-Interactive Interface and Human Authorship

3.1 Pseudo-Interactive UI System

The conceptual core is three game-style notification interfaces that reframe the film as a “system session.” Generated entirely in Nano Banana Pro from one template, they share a flat, high-contrast language: a 420×280 px rectangular frame (no rounded corners), 2 px accent border, pure-black background at 75–80% opacity, pixel-style typography, and a bottom progress bar with an action button. The three interfaces differ in when they appear, their accent color, and their function; Table 2 summarizes the three.

Interface #3’s statistics block reads “Total Time 06:00”, “Scenes 6/6”, “Completion 100%”, coldly summarizing a life that was never actually lived.

3.2 Human Intervention Points

Despite being 100% AI-generated, the film depends on human creative intervention at key decision points to keep the result from being mere tool output.

The protagonist’s face was first built through prompting yet was initially too neutral; working from that sample, the features were fine-tuned through further prompts and then locked as a three-view sheet with varied expressions, pulling “what she looks like” out of the model’s randomness and placing it under authorial control. Costumes were not invented from text: real clothing references were fused with the film’s visual-style keywords after a pure-prompt attempt gave weak texture and material quality, and the move from text to reference was an aesthetic judgment rather than a tool default.

Each scene was generated on its own rather than extended from one anchor; the resulting chromatic, costume, and temporal fractures were left intentional, an authorial choice rather than a tool defect. For the subway searching shot, one base single-shot scene image was first generated and then decomposed through prompting

into multiple angles of the same space, with the spatial layout fixed before generation so the tool produced the angles coherently.

Gemini assisted with the script and storyboard, but I refined the scene order, the emotional arc, and the Cantonese colloquial register, and I also designed the “FORCE EXIT / external interrupt detected” popups that bridge the fictional narrative and its medium.

The sound design carries further authorial decisions. A bed of white noise and city-road traffic — cars passing in the distance — underlies every scene, so the world always sounds urban and inhabited; the subway-station segment adds a rapidly passing train and, in its malfunction portion, an electric-current leakage buzz; the final office return layers a busy workplace ambience with telephone rings, so the protagonist is plainly not alone. None of these sound fields answer a cursor, and that non-response is the point.

4 Staging the Illusion of Agency

The pseudo-interactive form is not a technical compromise but the film’s central argument. The virtual world on screen is not generated live from viewer input; it is a pre-authored, world-model-style closed space. The seemingly interactive interface is only a formal shell — I fixed its branches in advance. The form is ‘pseudo’ because the on-screen interactivity is simulated, not real: the cursor and interfaces respond to no input, and the film plays identically whether or not a viewer interacts. That manufactured agency is the argument.

At the content level, the film injects the social and affective texture of the real world into this virtual space: everyday locations such as the convenience store, the subway, and the office, together with the blending of real-world garments and synthetic faces, weave the grain of real life into virtual imagery. At the authorial level, my subjective consciousness — the lived sense of alienated routine and the awakening that one is “still living” — is refracted through the toolchain, mediated and translated, into the virtual world. This reading builds on Huang et al.’s human-LLM co-creation, in which real footage is refracted into a fabricated narrative [3].

Pseudo-interaction thus becomes a mirror: the viewer believes they are choosing, yet what they see reflected is the predetermined, false agency within algorithmic and systemic world models.

The screen presents a world-model-like surface—a seemingly coherent and strictly regulated space that appears to predict the next move. No system can deduce the state of the world, no simulation runs behind the scenes, and the cursor movement does not revise any underlying representation. What the audience sees is an artificially constructed world that does not respond to anything, and the distance between this appearance and a truly responsive model is precisely the chasm the film wants us to experience.

Contemporaneous work moves in the opposite direction. TOUCH-STAR (Chen et al. [4]), for example, is an immersive VR interactive film in which gesture and voice input trigger real-time generative-AI content — environment textures, sound, 3D models, and dialogue — that reshapes the narrative as the viewer acts. This film works on a single 2D screen and stages the absence of that generation: the interface responds to nothing, and the world behind it is fixed. The contrast is precisely the point — a genuine generative-interactive

system (built for headset-based VR) and a pseudo-interactive one (built for a flat, pre-authored frame) evoke the same feeling of choice while differing on whether any choice is ever actually made.

The work asks its viewer for nothing beyond sitting through a closed system; a second viewing turns that passivity into something more active, as every cursor hover, every loading glitch, and every “FORCE EXIT” popup retroactively marks the medium itself. The form is also suited to AIGC specifically because the virtual world can be authored once and held static across the runtime; what changes is not the labor cost of building that world, but who is allowed to refuse the viewer a real choice. Pseudo-interactive AIGC film thus names a position rather than a workflow: the director keeps every branch closed, and the audience is asked only to recognize the closing.

5 Copyright Declaration

I confirm that this submission is an original work and that all materials included in the film have been created by me or are used with appropriate permissions. By submitting this work, I grant SIGGRAPH Asia 2026 and the workshop organizers a non-exclusive, royalty-free license to screen the film during SIGGRAPH Asia 2026 and to use the submitted materials for future workshop-related activities, archival purposes, and promotional or publicity materials. I retain the copyright of my work and assume full responsibility for any copyright-related issues arising from this submission.

Asset Ownership Details:

Visual Assets: 100% AI-generated (using Midjourney v7, Nano Banana Pro, Seedance 2.0 within Tapnow platform); all outputs are original content created by me through prompt engineering, not reproductions of existing copyrighted works.

Audio Assets: Dialogue voice synthesized via Seedance 2.0; Sound effects sourced from royalty-free libraries (properly licensed); No copyrighted music samples used.

Interface Design: Original design created by me; Any AI-assisted UI sketches generated via Nano Banana Pro are original compositions, not copies of existing interfaces.

Script/Story: Original script and narrative concept fully developed by me; AI tools (Gemini) used only for ideation assistance, not final content generation.

References

- [1] Y. Atzmon, R. Gal, Y. Tewel, Y. Kasten, and G. Chechik. Motion by Queries: Identity-Motion Trade-offs in Text-to-Video Generation. *arXiv preprint arXiv:2412.07750*, 2025.
- [2] M. Zheng et al. VideoGen-of-Thought: Step-by-step generating multi-shot video with minimal manual intervention. *arXiv preprint arXiv:2412.02259*, 2025.
- [3] X. Huang, X. Zhong, Z. Zhong, T. Papatheodorou, and K. Yip. Cinema Meowdiso: Films Co-Created by Human, Large Language Models and a Cat. In *Proceedings of SIGGRAPH Asia 2024 (Art Paper)*. ACM, 2024. 10.1145/3680530.3695443
- [4] Y. Chen, Y. Hu, and X. Shen. “TOUCHSTAR”: Exploring Generative AI-Powered Interactive Narrative in Cinematic Virtual Reality. In *Proceedings of SIGGRAPH Asia 2025 (XR)*. ACM, 2025. 10.1145/3761667.3761945
- [5] Y. Li, H. Shi, B. Hu, L. Wang, J. Zhu, J. Xu, Z. Zhao, and M. Zhang. Anim-Director: A Large Multimodal Model Powered Agent for Controllable Animation Video Generation. In *Proceedings of SIGGRAPH Asia 2024 (Conference Papers)*. ACM, 2024. 10.1145/3680528.3687688
- [6] H. Shi, Y. Li, X. Chen, L. Wang, B. Hu, and M. Zhang. AniMaker: Multi-Agent Animated Storytelling with MCTS-Driven Clip Generation. In *Proceedings of SIGGRAPH Asia 2025 (Conference Papers)*. ACM, 2025. 10.1145/3757377.3764009