

“But”: Pseudo-Interactive AIGC Film

Exploring the Illusion of Agency through Simulated Interaction and Game-Style UI

Shiqin Hou

HKUST

shouaf@connect.ust.hk

Abstract

“But” is a 6-minute AI-generated narrative short that treats the illusion of choice as its central subject. Within an integrated AI platform (Tapnow), a persistent on-screen cursor and three game-style notification interfaces frame the film as a predetermined “system session”: the viewer is invited to act, yet every branch is authored in advance. This technical report documents the end-to-end pipeline realized through a pseudo-interactive AIGC filmmaking framework (Gemini for scripting and stepwise shot planning, Midjourney v7 for scene stylization, Nano Banana Pro for character and UI assets under a unified aesthetic, Seedance 2.0 for animation, CapCut for compositing), the methods used to hold a single protagonist coherent across six costume changes, and the design of the pseudo-interactive UI as a formal shell for the film’s argument about false agency. This report discusses the work as a case of human-led, tool-mediated authorship in AIGC filmmaking.

CCS Concepts

• Computing methodologies → Computer graphics; • Applied computing → Media arts.

Keywords

AI-generated film, pseudo-interactive, game UI integration, character consistency, AIGC workflow, human-AI co-creation

ACM Reference Format:

Shiqin Hou. 2026. “But”: Pseudo-Interactive AIGC Film: Exploring the Illusion of Agency through Simulated Interaction and Game-Style UI. In *SIGGRAPH Asia 2026 AI Film Frontiers Workshop (SA '26)*. ACM, New York, NY, USA, 3 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

But is a 6-minute AIGC pseudo-interactive narrative short. This report proposes a pseudo-interactive AIGC filmmaking framework – a single authoring method that orchestrates off-the-shelf generative tools into one pre-authored, world-model-style film space wrapped in a game-style interface. The work grew out of a lived contradiction of wage labor: daily life had hardened into a script I only performed, until a glimpse of open sky from a pantry room

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
SA '26, Kuala Lumpur, Malaysia

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-X-XXXXXX-XX-X/26/12
<https://doi.org/XXXXXXX.XXXXXXX>

Table 1: The framework instantiated across tools: each of the six design stages with its authorial method and realizing tool, plus the video engine (Seedance 2.0) used for scene motion.

Design stage	Authorial method	Tool
Scene-first	Script + shot planning	Gemini
Character assets	Locked face, 3-view sheet	Nano Banana Pro
Real-garment	Real outfits + style (real refs)	Midjourney v7 + real refs
Fragmentation	Intentional fractures	authorial
Multi-view	Base → multi-angle	Nano Banana Pro
UI shell	UI generated, then cut into the final edit	Nano Banana Pro + CapCut
Scene motion	Animates each generated scene	Seedance 2.0

broke it – a fresh awareness of observing the self rather than experiencing it firsthand. Atzmon et al. [1] name the identity–motion trade-off that makes a consistent protagonist hard, and Zheng et al. [2] propose step-by-step, minimal-intervention generation – together they point to consistency and planning as the real engineering problems, not the models; Li et al. [5] and Shi et al. [6] orchestrate cross-scene consistency, while Huang et al. [3] show how human–LLM co-creation turns real footage into a fabricated narrative. The film takes the one move none of them make: where Chen et al. [4] drive a cinematic VR narrative through real-time, input-triggered generative-AI interaction, *But* inverts that premise and simulates interaction from a fixed script, turning the viewer’s sense of control into a media-art claim about living inside systems that feel free yet hand us only the appearance of choice. Carrying this out meant working entirely within commodity tools – Gemini, Nano Banana Pro, Midjourney v7, Seedance 2.0, and CapCut (Table 1) – yet two difficulties arose: keeping one protagonist legible across six costumes and discontinuous time, and unfolding the pseudo-interaction so the viewer’s sense of control reads as real before it is withdrawn. The response was to build a small system rather than patch scenes: a reusable character-asset library, and a structure that withholds the film’s fantastical nature until the final moments, so the viewer first inhabits the daydream as real and is then retroactively shown it was simulation. The result is a 6-minute MP4 video (H.264, 2560×1440, 30 fps, Cantonese with English subtitles).

2 The Pseudo-Interactive AIGC Framework

The generative stages run within **Tapnow**, which consolidates the model tools – Gemini, Midjourney v7, Nano Banana Pro, and Seedance 2.0 – into one workspace; the generated clips and UI images are then assembled in CapCut. The six design stages map onto these tools as Table 1 shows.

The framework is a single production chain in which those six stages are orchestrated across commodity tools (Table 1). Subsections 2.1–2.3 realize design stages (2)–(5); the shell (stage 6) is the subject of Subsection 3.1.

2.1 Multi-Appearance Character Consistency

Design stage (2): creating reusable assets and building a character-asset library.

The first major difficulty is keeping one recognizable protagonist stable across six costumes, six lighting setups, and six scenes. Text-to-video generators let a single prompt drift freely, so the same face arrives at different generations with different features or styling – the identity–motion trade-off Atzmon et al. [1] describe. Anim-Director [5] and AniMaker [6] orchestrate cross-scene consistency through autonomous agents; both are bespoke research systems, not the off-the-shelf tools used here. This film borrows only the paradigm – that consistency must be engineered rather than left to the generator – and removes the drift at the source. The face was produced by prompting yet was initially too neutral; working from that sample, the features were fine-tuned through further prompts and then locked as a three-view sheet with varied expressions that becomes a reusable asset fed into costume design, scene fusion, and video generation. Costumes are not invented from text: real clothing references are fused with the film’s style keywords, replacing an earlier pure-prompt attempt whose texture and material quality fell short. The result is that audiences identify the same person through every costume change, which lets the final reveal achieve its maximum impact.

The deliberate chromatic, costume, and temporal fractures between scenes are not defects but an authorial choice: each scene was generated independently, and the resulting discontinuity carries the film’s theme of a pluralistic life (Subsection 2.3). The recurrence of one face is also what lets the film read as a single person’s routine rather than six unrelated figures: the same features survive every costume and lighting change, so the viewer assembles one continuous subject out of clips that were never shot together – the week repeats a protagonist, not a story. The sameness is authored, not observed: a stable, repeatable persona is delivered as given.

2.2 Multi-Camera Spatial Continuity

Design stage (5): from a single plane to multiple three-dimensional views.

Creating the subway-station scene raised a spatial-continuity problem, and the pseudo-interactive element made that problem more acute. The design required a first-person POV shot in which the camera appears to search for the protagonist; after she walks off-frame, two empty shots play under electromagnetic-interference sound, a “FORCE EXIT” popup explains the interruption, and once the cursor is clicked the camera returns and finds her leaning against a wall, asking “Are you looking for me?”. Holding this POV insertion continuous with the surrounding special-angle shots – foreground overhead, side tracking, low-angle rotation – demands that every angle describe the *same* space. Rather than draw a 2D floor plan, the method anchors all angles to one generated base: a single-shot image of the intended location is produced first, then decomposed through prompting into multiple camera angles of that same spot, so

geometry stays coherent across the cut. The mouse cursor persists across the empty shots and the return, marking the screen-mediated nature of the experience; a diegetic sound field – a passing train and the electric leakage of the glitch – implies off-screen movement even while the protagonist is absent.

Anchoring every angle to a single generated base keeps geometry coherent, but it also means the ‘multiple camera views’ are recombinations of one authored image rather than independent captures. That constraint is the point: the sequence presents the illusion of a navigable space while every vantage is in fact derived from the same fixed frame. The coherence the viewer reads as spatial continuity is therefore pre-authored closure, the same closure the film’s world-model reading names as a closed, non-responsive system.

2.3 Fantasy and Reality Controllability

Design stage (3) and (4): fusing real garments and fragmenting scenes on purpose.

The contradictory narrative demand is that audiences take the fantasy sequences as real until the final reveal, yet catch on rewatching the clues that expose the artifice. The film meets this through deliberate per-scene color design rather than a binary color code. Each scene carries a consciously chosen dominant tone – the store’s black-and-white checkerboard, the bedroom’s cool-and-warm mix, the restaurant’s yellow, the stairwell’s cement gray, the subway’s red, the rooftop’s night-sky blue, and the office’s deep blue – yet none of these tones is assigned to mark a scene as fantasy or as real. The fantasy/reality boundary is instead carried by the protagonist’s seemingly arbitrary, narratively unmotivated costume changes, which blur imagined and real and suggest every scene belongs to one unstable mental space, and by scene transitions that jump illogically between day and night: a temporal disorientation that gives the experience its dreamlike, constructed feel. Two “system loading” moments use brief glitch effects that appear at first to be stylistic, then reveal themselves as diegetic hints about the medium. The three game interfaces (Subsection 3.1) add a further artificial-signal layer, retrospectively reframing the whole fantasy portion as something occurring inside a digital frame.

2.4 Sound, Scene Order, and Passivity

Design stage (6): the pseudo-interactive UI shell and the post-production layer that binds the work.

Editing and assembly are the final creative act: they make the pseudo-interactive reading legible. The scenes – convenience store, bedroom, restaurant, stairwell, subway, rooftop birthday party, office – are cumulative, not progressive: the film loops through a week’s worth of work, the rooftop birthday party being one segment in the loop, not a release. The form is pseudo-interactive, but the routine is real, and the sound and scene sequence emphasize the gap between the two. The sequence of scenes, the timing of each glitch, the instant a pop-up appears – these determine how the work is read: the viewer is watching a session, not a story being told. Design stage (6) refers to the post-production layer: the argument resides not in any single shot but in the blank spaces chosen during editing.

Table 2: The three pseudo-interactive interfaces.

#	Trigger	Accent	Function
1	Opening convenience-store scene	white/cyan	Session start
2	Subway scene (protagonist off-frame)	red	Warning message
3	Post-reveal hard cut	white/green	Statistics

Two other production decisions, in sound direction and narrative pacing, also reinforce the pseudo-interactive argument. The narrative soundscape, deliberately added in post-production, reflects a denial of the viewer’s autonomy, much like the visual style. Every scene is laid over a bed of white noise and city-road traffic — cars passing in the distance — so the world always sounds urban and inhabited; the subway adds a speeding train and the hum of leaking electricity. But none of these soundscapes respond to cursor clicks. The “FORCE EXIT” popups, system loading failures, and the ever-present mouse cursor appear as the final layer of this emphasis: the interface symbolizes autonomy, while the world rejects it.

3 Pseudo-Interactive Interface and Human Authorship

3.1 Pseudo-Interactive UI System

The conceptual core is three game-style notification interfaces that reframe the film as a “system session.” Generated entirely in Nano Banana Pro from one template, they share a flat, high-contrast language: a 420×280 px rectangular frame, 2 px accent border, pure-black background, pixel-style typography, and a bottom progress bar with an action button. Table 2 summarizes the three.

Interface #3’s statistics block reads “Total Time 06:00”, “Scenes 6/6”, “Completion 100%”, coldly summarizing a life that was never actually lived.

3.2 Human Intervention Points

Despite being 100% AI-generated, the film depends on human creative intervention at key decision points to keep the result from being mere tool output.

Where the tool offered a default, the authorial move was to refuse it: the face was locked as a three-view sheet rather than left to drift, costumes were fused from real references rather than invented from text, and each scene was generated on its own rather than extended from one anchor, leaving the resulting fractures intentional. Gemini assisted with the script and storyboard, but I refined the scene order, the emotional arc, and the Cantonese colloquial register, and I also designed the “FORCE EXIT / external interrupt detected” popups that bridge the fictional narrative and its medium. The sound design carries the same logic: none of its fields answer a cursor, and that non-response is the point.

4 Staging the Illusion of Agency

The pseudo-interactive form is not a technical compromise but the film’s central argument. The virtual world on screen is not generated live from viewer input but is a pre-authored, world-model-style closed space, and the seemingly interactive interface is

only a formal shell — I fixed its branches in advance. The form is ‘pseudo’ because the on-screen interactivity is simulated, not real: the cursor and interfaces respond to no input, and the film plays identically whether or not a viewer interacts.

At the content level, the film injects the social and affective texture of the real world into this virtual space: everyday locations such as the convenience store, the subway, and the office, together with the blending of real-world garments and synthetic faces, weave the grain of real life into virtual imagery. At the authorial level, my subjective consciousness — the lived sense of alienated routine and the awakening that one is “still living” — is refracted through the toolchain into the virtual world. Pseudo-interaction thus becomes a mirror: the viewer believes they are choosing, yet what they see reflected is the predetermined, false agency within algorithmic and systemic world models.

The form suits AIGC because the virtual world can be authored once and held static across the runtime; what changes is not the labor cost of building that world, but who is allowed to refuse the viewer a real choice. Pseudo-interactive AIGC film thus names a position rather than a workflow: the director keeps every branch closed, and the audience is asked only to recognize the closing.

5 Copyright Declaration

This submission is an original work; all film materials were created by me or are used with appropriate permissions. By submitting, I grant SIGGRAPH Asia 2026 and the workshop organizers a non-exclusive, royalty-free license to screen the film and to use the materials for workshop, archival, and promotional purposes. I retain copyright and assume full responsibility for any copyright-related issues.

Visual and Audio Assets: 100% AI-generated (Midjourney v7, Nano Banana Pro, Seedance 2.0; dialogue synthesized via Seedance 2.0; sound effects from licensed royalty-free libraries); all outputs are original content, not reproductions of copyrighted works.

Interface Design and Script/Story: Original designs and narrative concept developed by me; AI-assisted UI sketches (Nano Banana Pro) are original compositions, and AI tools (Gemini) were used only for ideation, not final content generation.

References

- [1] Y. Atzmon, R. Gal, Y. Tewel, Y. Kasten, and G. Chechik. Motion by Queries: Identity-Motion Trade-offs in Text-to-Video Generation. *arXiv preprint arXiv:2412.07750*, 2025.
- [2] M. Zheng et al. VideoGen-of-Thought: Step-by-step generating multi-shot video with minimal manual intervention. *arXiv preprint arXiv:2412.02259*, 2025.
- [3] X. Huang, X. Zhong, Z. Zhong, T. Papatheodorou, and K. Yip. Cinema Meowdiso: Films Co-Created by Human, Large Language Models and a Cat. In *Proceedings of SIGGRAPH Asia 2024 (Art Paper)*. ACM, 2024. 10.1145/3680530.3695443
- [4] Y. Chen, Y. Hu, and X. Shen. “TOUCHSTAR”: Exploring Generative AI-Powered Interactive Narrative in Cinematic Virtual Reality. In *Proceedings of SIGGRAPH Asia 2025 (XR)*. ACM, 2025. 10.1145/3761667.3761945
- [5] Y. Li, H. Shi, B. Hu, L. Wang, J. Zhu, J. Xu, Z. Zhao, and M. Zhang. Anim-Director: A Large Multimodal Model Powered Agent for Controllable Animation Video Generation. In *Proceedings of SIGGRAPH Asia 2024 (Conference Papers)*. ACM, 2024. 10.1145/3680528.3687688
- [6] H. Shi, Y. Li, X. Chen, L. Wang, B. Hu, and M. Zhang. AniMaker: Multi-Agent Animated Storytelling with MCTS-Driven Clip Generation. In *Proceedings of SIGGRAPH Asia 2025 (Conference Papers)*. ACM, 2025. 10.1145/3757377.3764009